


Gradient Descent – Super Simple and Detailed Explanation (in English)

What is Gradient Descent?

Gradient Descent is a method to find the minimum of a function.

 In machine learning, our goal is to reduce the error or loss function (how wrong the model is).

We take small steps in the direction where the error decreases most quickly — that direction is called the negative gradient.

Real-Life Analogy (Visual Intuition)

Imagine you're standing on a mountain in the dark:

- You can't see
- You can only feel the slope under your feet
- You want to reach the lowest valley
- So you take small steps downhill

That's exactly what Gradient Descent does — but instead of a hill, it's a math function.

Basic Setup – Using Simple Variables

Let's say you have a simple linear regression model:

$$\hat{y} = m * x + b$$

Where:

- x = input
- \hat{y} = predicted output
- y = actual output
- m = slope (weight)
- b = intercept (bias)

Loss Function (Error Measurement)

We want to measure how far off our prediction is. The most common formula:

$$\begin{aligned} J(m, b) &= (1/n) * \sum (y - \hat{y})^2 \\ &= (1/n) * \sum (y - (mx + b))^2 \end{aligned}$$

This is called the Mean Squared Error (MSE).

$J(m, b)$ tells us how bad our model is, depending on m and b .

Gradient Descent Steps

We want to update m and b in a way that reduces the error $J(m, b)$.

Gradient Descent Update Rules

We update our weights like this:

$$\begin{aligned} b &:= b - \alpha * \partial J / \partial b \\ m &:= m - \alpha * \partial J / \partial m \end{aligned}$$

Where:

- α = learning rate (how big a step we take)
- $\partial J / \partial b$ = derivative of cost function with respect to b
- $\partial J / \partial m$ = derivative of cost function with respect to m

Partial Derivatives (With Simple Math)

Using calculus to find derivatives of the loss function:

$$J(m, b) = (1/n) * \sum (y - (mx + b))^2$$

Then:

$$\begin{aligned} \partial J / \partial b &= (-2/n) * \sum (y - \hat{y}) \\ \partial J / \partial m &= (-2/n) * \sum x * (y - \hat{y}) \end{aligned}$$

Final Update Rules (Simplified):

$$\begin{aligned} b &:= b + \alpha * (2/n) * \sum (y - \hat{y}) \\ m &:= m + \alpha * (2/n) * \sum x * (y - \hat{y}) \end{aligned}$$

Note: The $+$ comes from the double negative ($-$ in rule and $-$ in gradient)

Step-by-Step Algorithm (Linear Regression)

1. Start with random values: $m = 0, b = 0$
2. For each step:
 - Compute predictions: $\hat{y} = mx + b$
 - Compute error: $y - \hat{y}$
 - Compute derivatives
 - Update m and b using formulas
3. Repeat until the error becomes very small (i.e., converges)

Learning Rate (α) — Important!

- Too small: very slow learning
- Too large: model might jump over the minimum and never settle
- Should be tuned carefully

Summary in Simple Words

Concept	Meaning
-----	-----
Gradient	Direction of steepest slope
Loss Function	Measures model error
Learning Rate (α)	Step size for updates
Update Rule	Formula to update m and b
Goal	Minimize loss function by adjusting model weights